

Model Compression

영남대학교
차세대 컴퓨터 시스템 연구실
석사과정생 이원호

Model compression and hardware acceleration for neural networks: A comprehensive survey.

Proceedings of the IEEE, 2020

차세대 컴퓨터 시스템 연구실

이원호

Basic Model Architecture

- MLP (Multi-Layer Perceptron)
 - Feedforward Fully Connected (FC)
 - Efficient processing 1-dimensional features
- CNN (Convolutional Neural Network)
 - Efficient processing 2-dimensional features
- RNN (Recurrent Neural Network)
 - Efficient processing sequence features
 - Handling temporal dynamics

Basic Model Architecture

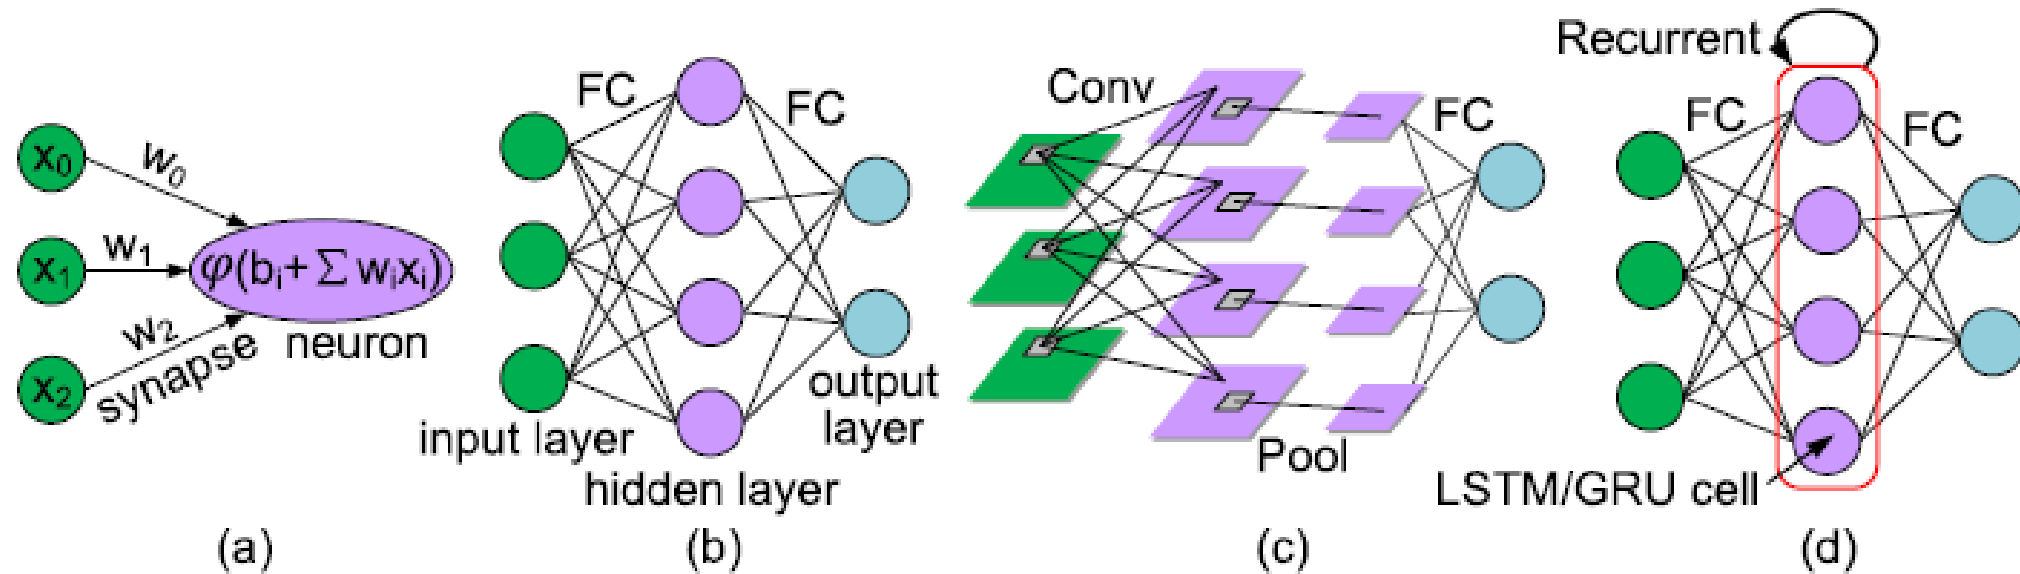


Fig. 2. *Neural network structures. (a) Single neuron. (b) MLP. (c) CNN. (d) RNN.*

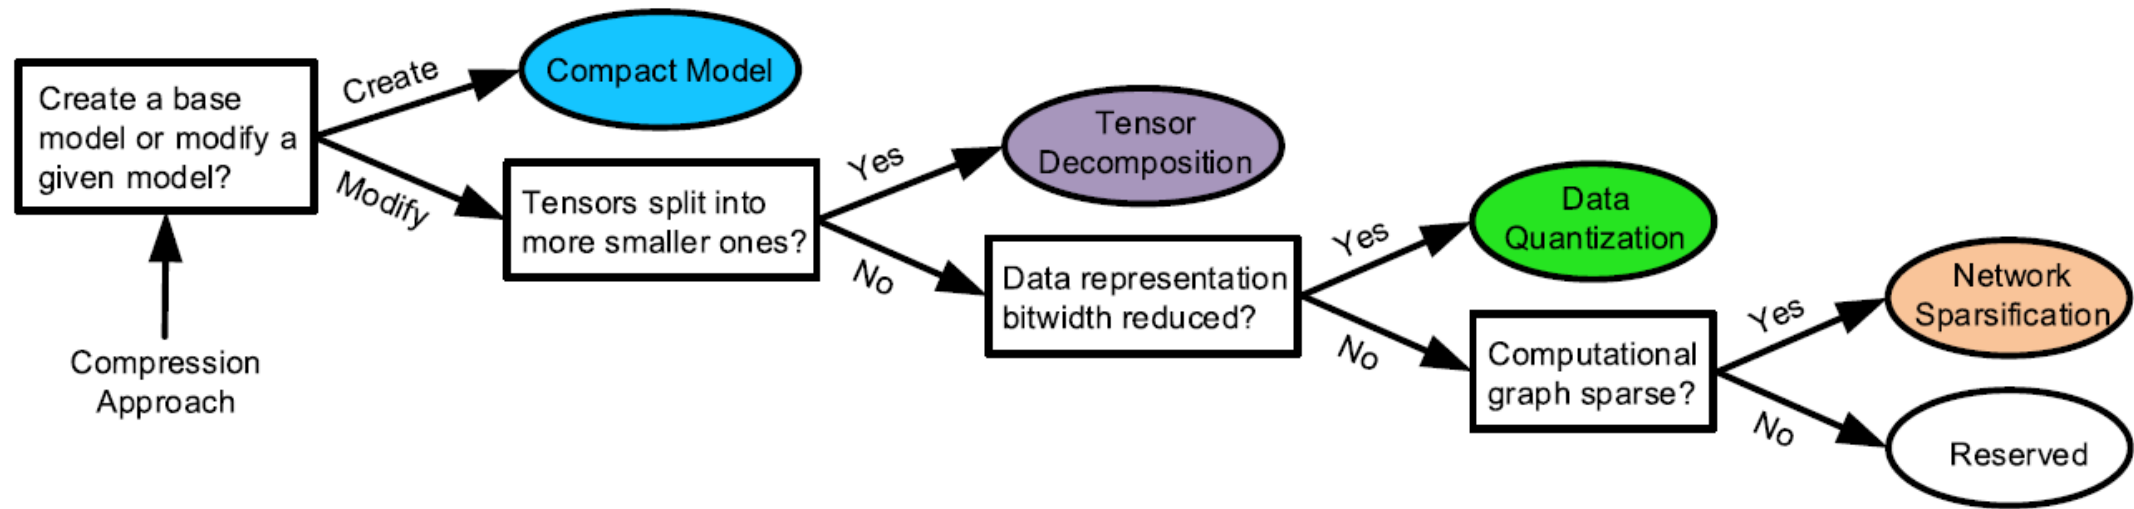
Compression Approaches

- Compact model
 - Design smaller base models that can still achieve acceptable application accuracy
- Tensor decomposition
 - Decompose the bloated parameters into a series of smaller matrices or tensors
- Data Quantization
 - Reduce the number of data bits
- Network sparsification
 - Reduce the connections/neurons to compress the original model

Compression Approaches

- **Compact model**
 - Design smaller base models that can still achieve acceptable application accuracy
- **Tensor decomposition**
 - Decompose the bloated parameters into a series of smaller matrices or tensors
- Data Quantization
 - Reduce the number of data bits
- Network sparsification
 - Reduce the connections/neurons to compress the original model

Taxonomy of neural network compression approaches



Compact Model

- Compact CNNs
- Compact RNNs
- Neural Architecture Search (NAS)

Compact CNNs

- CNN의 특징: 공유 가중치를 통한 지역 연결성 (Spatial correlation)

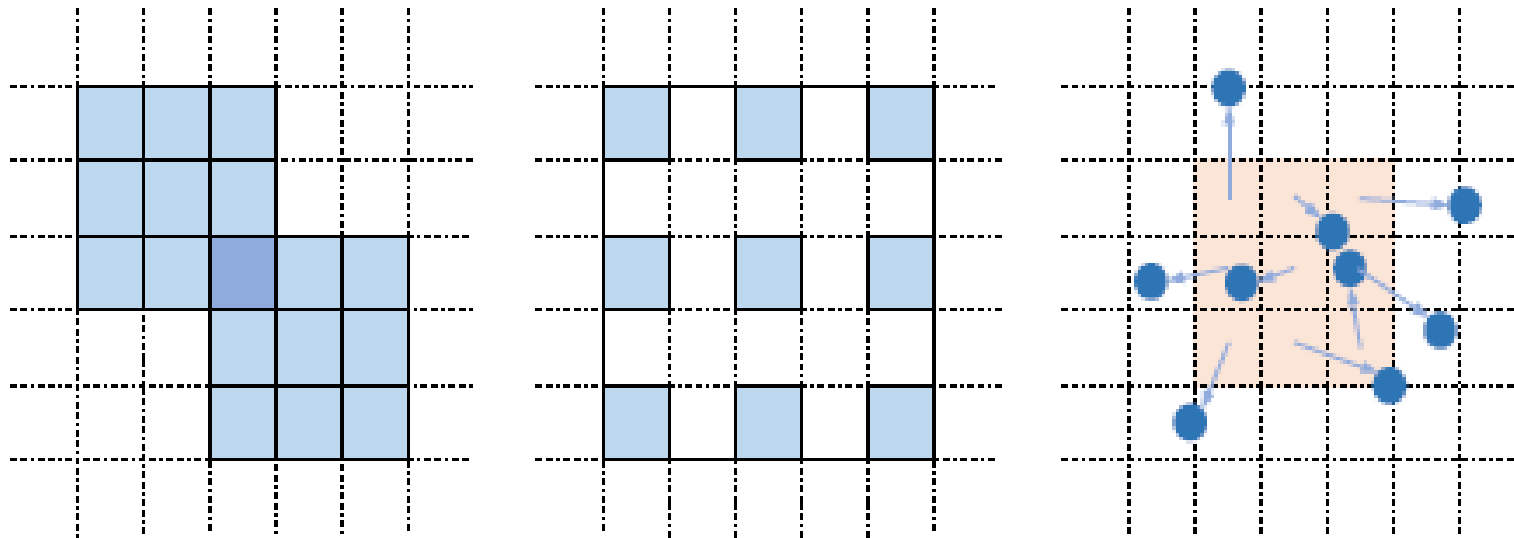


Fig. 5. (a) Sampling locations in standard Conv (stacked 3×3 conv). (b) and (c) Sampling locations in atrous Conv (dilated 3×3 conv and deformable 3×3 conv, resp.).

Compact CNNs

- Atrous Conv
 - Irregular kernels with holes dilations
- Deformable Conv
 - Generalization of atrous Conv
 - Add 2D offset factors

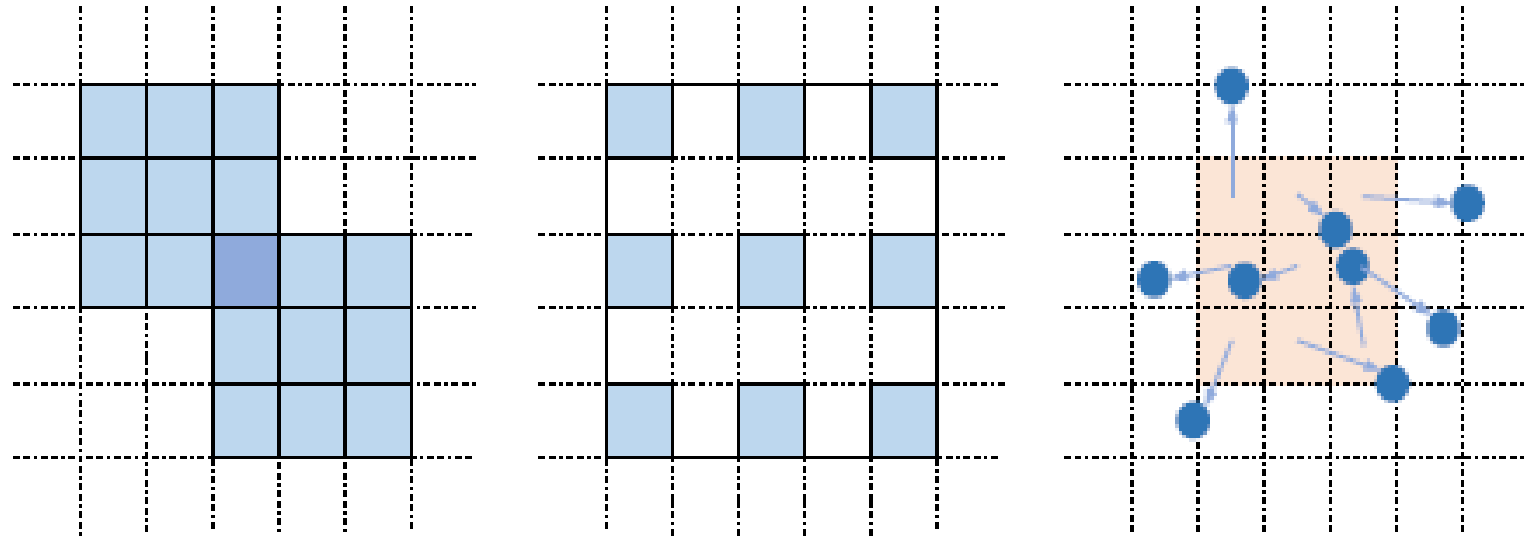


Fig. 5. (a) Sampling locations in standard Conv (stacked 3×3 conv). (b) and (c) Sampling locations in atrous Conv (dilated 3×3 conv and deformable 3×3 conv, resp.).

Compact CNNs

- 특징 맵이 다른 형태의 토폴로지로 연결될 때 생성 발생하는 관계
(interlayer channel correlation)

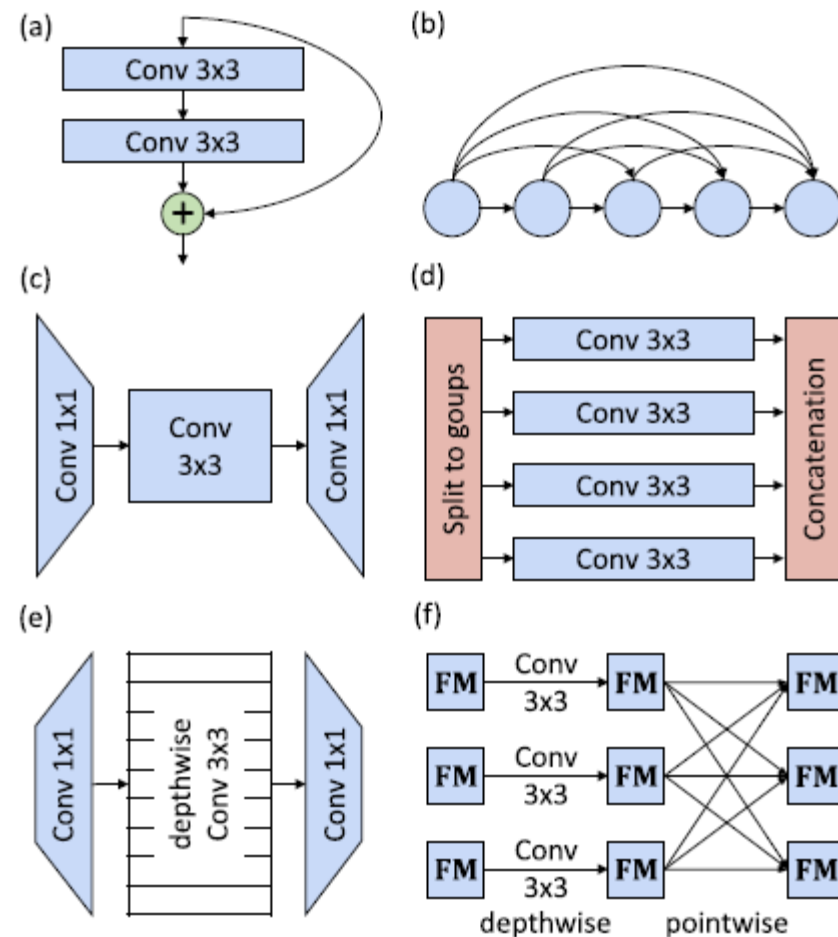


Fig. 6. Typical interlayer and intralayer channel correlations.
(a) Residual connection [30]. (b) Densely connected block [31].
(c) Bottleneck architecture [22]. (d) Group Conv [33]. (e) Reversed bottleneck architecture with depthwise Conv in between [3].
(f) Depthwise-separable Conv [34].

Compact RNNs (Unit Level)

- LSTM
 - Vanilla RNN에 비해 장기 의존성 향상
- GRU
 - Forget gate + input gate \rightarrow update gate
 - 불필요한 연결성 제거
- MGU
 - GRU에서 reset gate 와 update gate를 결합

Compact RNNs (Unit Level)

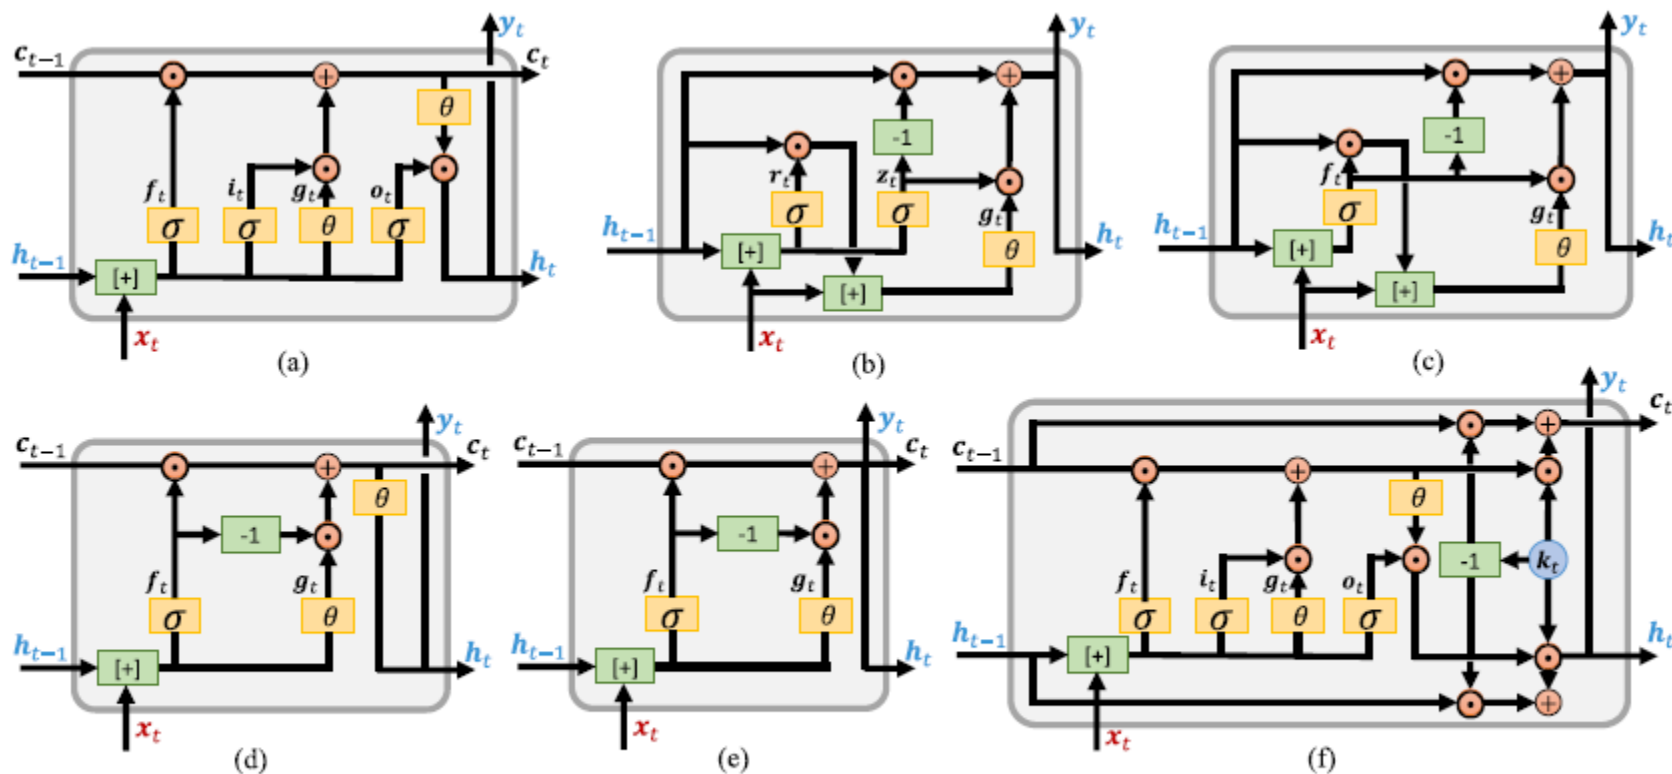


Fig. 7. Examples of standard and compact RNN units. (a) LSTM [18]. (b) GRU [19]. (c) MGU [40]. (d) S-LSTM [41]. (e) JANET [42]. (f) Phased LSTM [43].

Compact RNNs (Network Level)

- Linear recurrent projection
 - 차원의 hidden state 수를 줄임
- Factorized LSTM
 - LSTM의 weigh를 두 개의 작은 행렬로 축소시킴
- Residual connections
 - 여러 LSTM 간의 residual connection을 추가
- Skip-connection
 - 비연속적으로 연결된 RNN간에 직접적인 정보전달이 가능케함

Compact RNNs (Network Level)

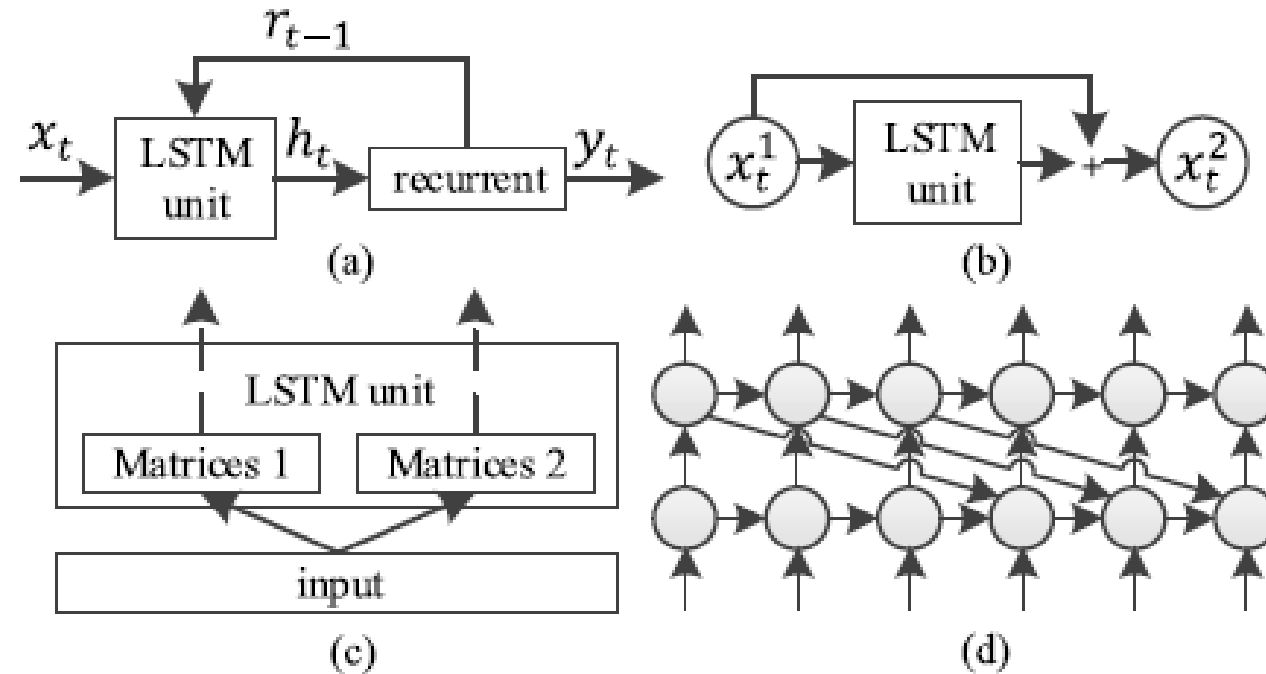


Fig. 8. Examples of compact RNNs in network level. (a) Linear recurrent LSTM [52]. (b) Residual LSTM [54]. (c) Factorized LSTM [53]. (d) Skip-connected RNN [55].

Neural Architecture Search

- 신경망의 종류가 다양해짐에 따라 적합한 구조를 찾는 작업의 비용이 증가함
- 이를 머신러닝 기술을 적용해 자동화하는 탐색 기술

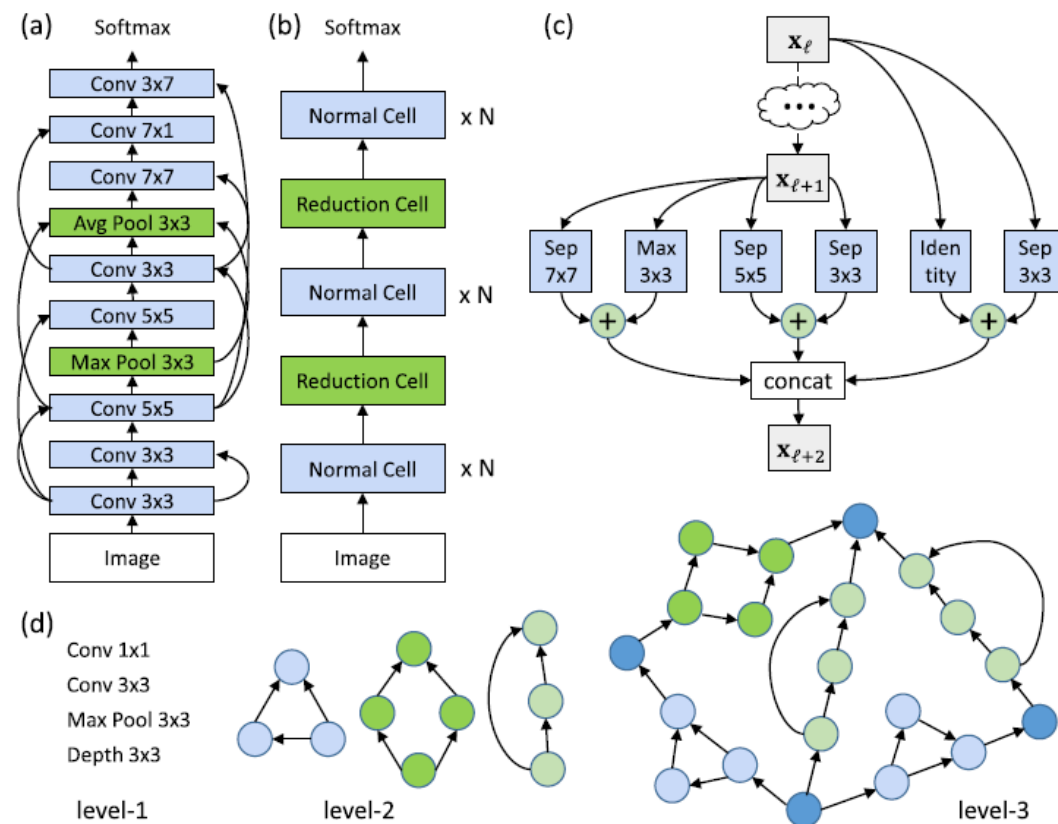


Fig. 9. Illustration of NAS. (a) Full-network search [45]. (b) Cell-based search [58]. (c) Example of learned cell structures [66]. (d) Three-level hierarchical search [67].

Tensor Decomposition

- Low-Rank Matrix Decomposition
- Tensorized Decomposition

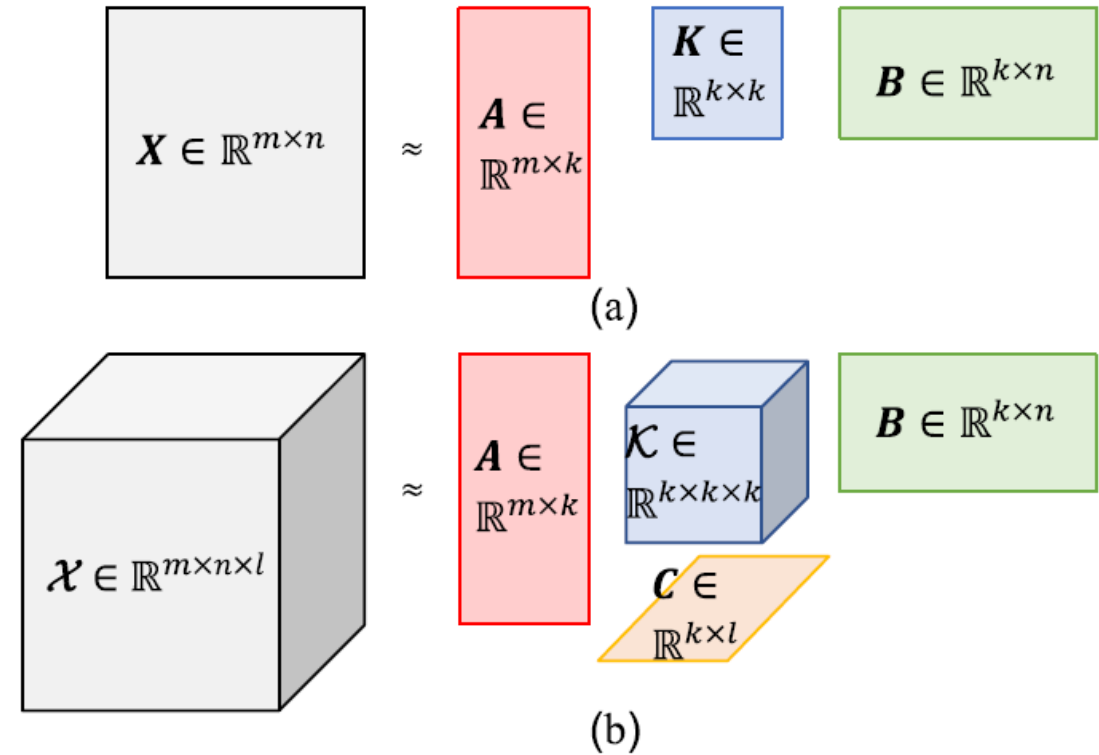


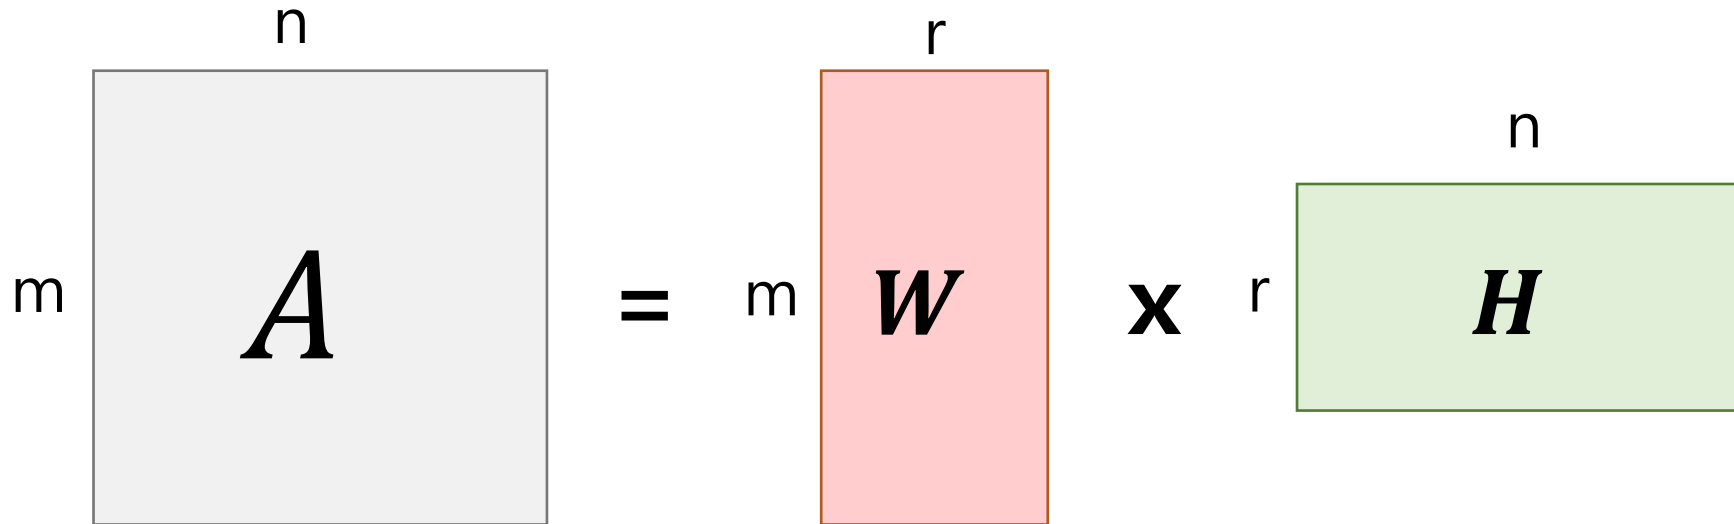
Fig. 11. Analogy between low-rank matrix decomposition and tensor decomposition. (a) Low-rank matrix decomposition with $k \times k$ kernel. (b) Low-rank tensor decomposition with $k \times k \times k$ kernel.

Low-Rank Matrix Decomposition

- Full-rank decomposition
- **Singular value decomposition (SVD)**

Full-rank decomposition

- 기존 A 행렬을 W 와 H 행렬의 곱으로 분해
- Spatial complexity : $O(mn) \rightarrow O(r(m+n))$



Singular value decomposition (SVD)

- 기존 A 행렬을 직교 행렬 U, V 와 대각행렬 S 로 분할
- Spartial complexity : $O(mn) \rightarrow O(r(m+n+1))$
- 복원을 위해 r 보다 작은 특정 값을 통해 기존 벡터가 가지는 특징값을 추출할 수 있음 ([예시 링크](#))

$$\begin{matrix} & n \\ m & \boxed{A} \end{matrix} = \begin{matrix} & r \\ m & \boxed{U} \end{matrix} \times \begin{matrix} & r \\ r & \boxed{S} \end{matrix} \times \begin{matrix} & n \\ r & \boxed{V^T} \end{matrix}$$

Tensorized decomposition

- Tucker and classical prolongation (CP)
- Tensorizing and curse of dimensionality
- HT(Hierarchical tucker) and dimension tree

Tucker and classical prolongation (CP)

- 행렬에서의 SVD를 텐서 차원으로의 단순 확장
- 같은 크기의 d차원 텐서
- $\chi \in \mathbb{R}^{n_1 \times n_2 \times \dots \times n_d}$, $K \in \mathbb{R}^{r_1 \times r_2 \times \dots \times r_d}$, $F^{(i)} \in \mathbb{R}^{r_i \times n_i}$
- $\chi \approx K \times F^{(1)} \times F^{(2)} \times \dots \times F^{(d)}$ 로 변환
- Spartial complexity : $O(n^d) \rightarrow O(nr + r^d)$

Tensorizing and curse of dimensionality

- 대규모 행렬 $W \in \mathbb{R}^{M \times N}$, $M = m_1 \times m_2 \times \cdots \times m_d$, $N = n_1 \times n_2 \times \cdots \times n_d$ 을 텐서 $\mathcal{W} \in \mathbb{R}^{m_1 n_1 \times m_2 n_2 \times \cdots \times m_d n_d}$ 로 변환 (tensorizing)
- 해당 텐서에 Tucker decomposition 적용시 공간 복잡도는 다음과 같다.
 - $O((mn)^d) \rightarrow O(dmnr + r^d)$ 여전히 지수차원의 공간 복잡도 존재
- 신경망에서 tensor의 표현은 행렬들의 연결 또는 저차원 텐서의 결합으로 표현되기 때문에, 이런 관계를 분할해 차원의 저주 해결 가능

HT(Hierarchical tucker) and dimension tree

- 차원의 총 개수 d 를 임의의 값 k 를 따라 $(1 \sim d)$ 를 $(1 \sim k), (1 \sim d-k)$ 차원의 두 개 텐서로 반복적으로 나눠 계산하는 방식을 HT라 함
 - 이 때 $k=1$ 을 기준으로 한다면 이를 TT(Tensor Train)이라 한다.

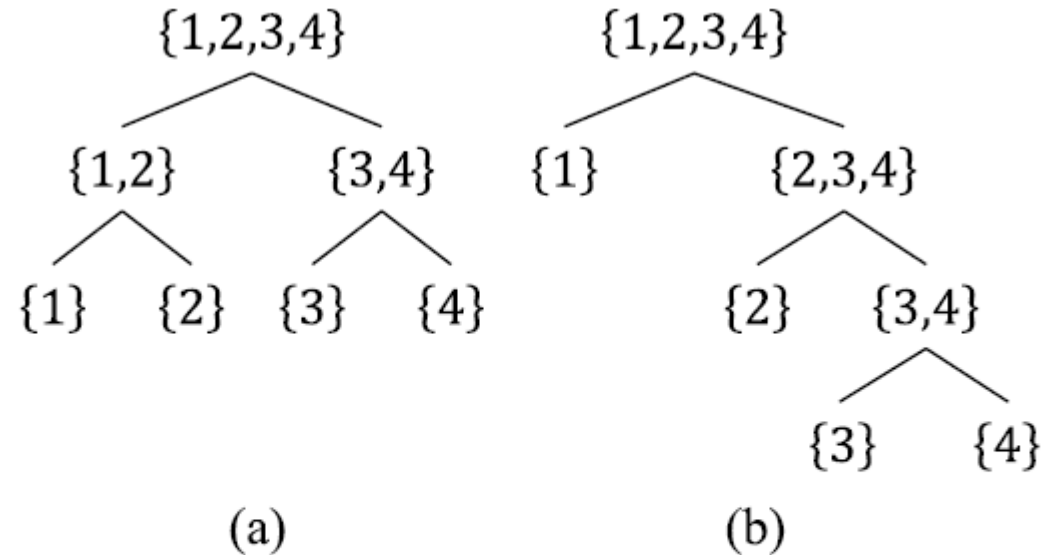


Fig. 12. Different dimension trees of tensor $\mathcal{X} \in \mathbb{R}^{n_1 \times n_2 \times n_3 \times n_4}$.
(a) Canonical dimension tree for the HT format. (b) Degenerate dimension tree for the TT format.

Reference

- [Model compression and hardware acceleration for neural networks: A comprehensive survey. Proceedings of the IEEE, 2020](#)
- [공돌이의 수학 정리 노트](#)

Model Compression

영남대학교

차세대 컴퓨터 시스템 연구실

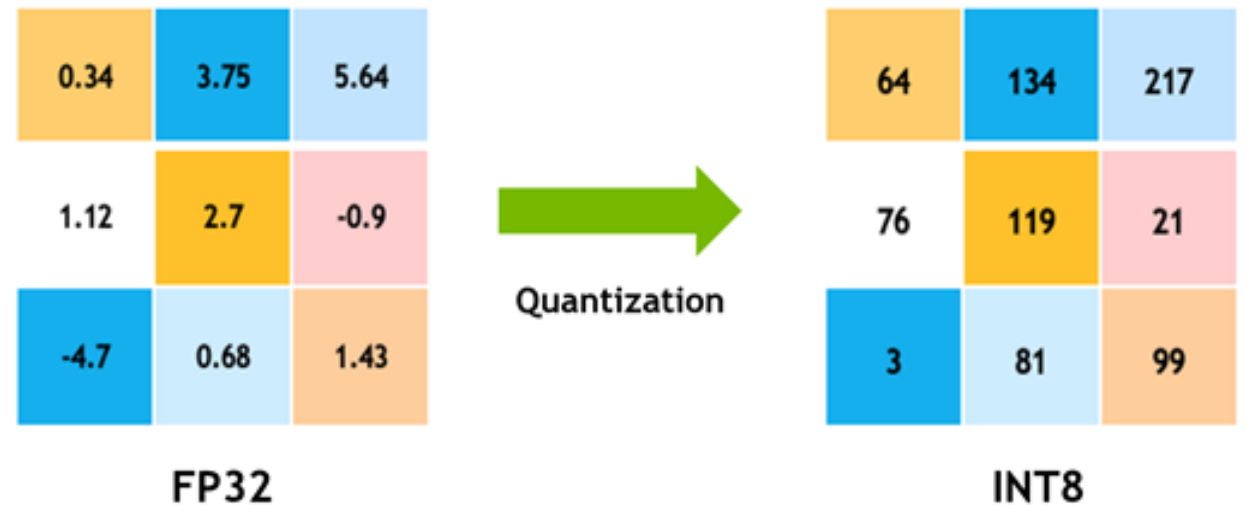
석사과정생 이원호

Compression Approaches

- Compact Model
 - Design smaller base models that can still achieve acceptable application accuracy
- Tensor Decomposition
 - Decompose the bloated parameters into a series of smaller matrices or tensors
- **Data Quantization**
 - Reduce the number of data bits
- **Network Sparsification**
 - Reduce the connections/neurons to compress the original model

Data Quantization

- Quantization Data Object
 - **What** to be quantized
- Quantization Considerations
 - **How** to quantize



Reduce bit-width → save the memory capacity

Quantization Data Object

- Three types of quantization target
 - **Parameter (Weight)**
 - Propagation data (Activation, Error)
 - Gradient and Update (Gradient, Weight Update)
- In recent studies, Quantization to BN layers is being discussed

Quantization Consideration

- Data Object
- Problem formulation
 - Heuristic / Optimization
- Distribution of discrete levels
 - Uniform / Non-uniform
- Level projection
 - Deterministic / Stochastic

Quantization Consideration

- Data Object
- Heuristic
 - Intuitively project data to discrete level
- **Problem formulation**
 - Heuristic / Optimization
- Distribution of discrete levels
 - Uniform / Non-uniform
- Level projection
 - Deterministic / Stochastic

$$Q(x) = \Delta \cdot \text{round}\left(\frac{x}{\Delta}\right)$$

- Optimization
 - Analytically approximation

$$\min_Q ||X - Q(X)||_2^2 \quad Q_i \in X_Q \text{ for all } i$$

Quantization Consideration

- Data Object
- Problem formulation
 - Heuristic / Optimization
- **Distribution of discrete levels**
 - Uniform / Non-uniform
- Level projection
 - Deterministic / Stochastic

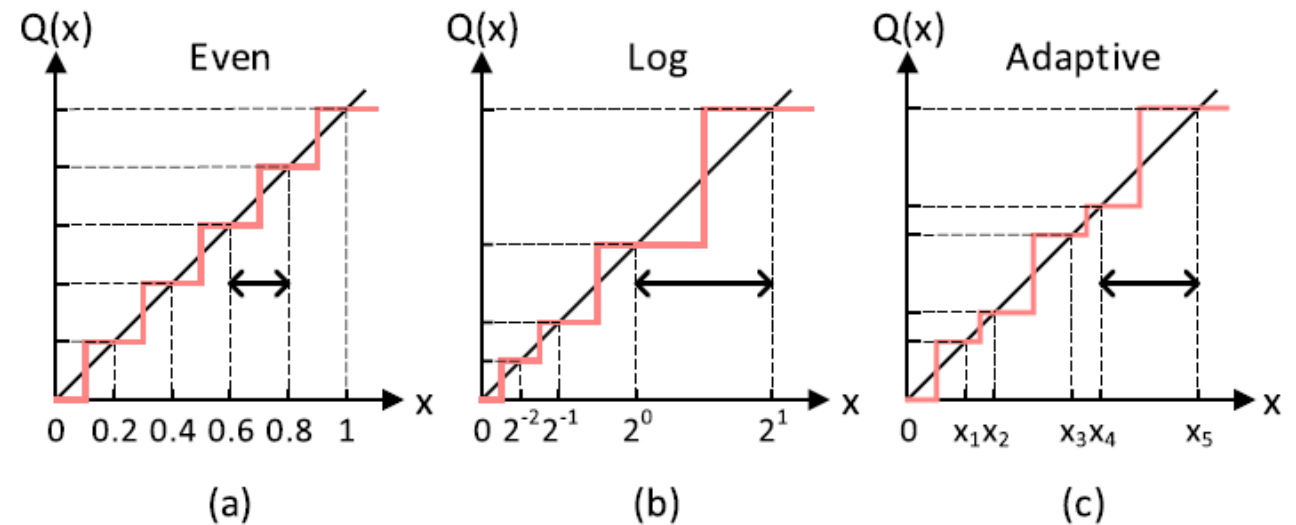
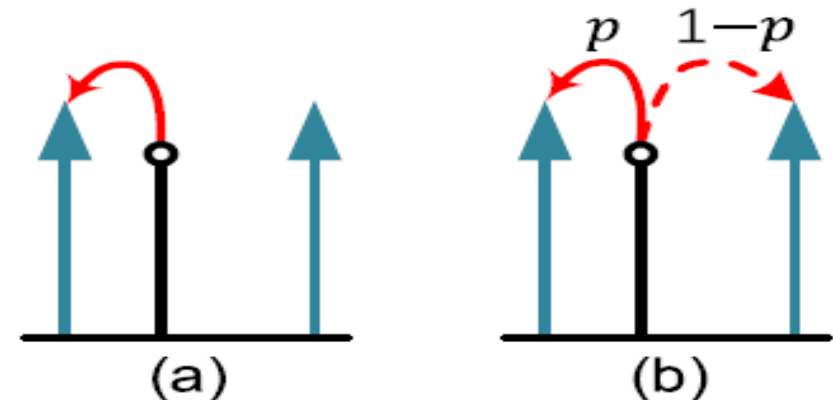


Fig. 15. *Distribution of discrete levels. (a) Uniform distribution with even step length. Nonuniform distribution with (b) logarithmic or (c) adaptive step length.*

Quantization Consideration

- Data Object
- Problem formulation
 - Heuristic / Optimization
- Distribution of discrete levels
 - Uniform / Non-uniform
- **Level projection**
 - Deterministic / Stochastic
- Deterministic way
 - Project data to nearest discrete level
- Stochastic way
 - Provides probabilities for two adjacent levels of data



Simple Quantization (single bit quantization)

- Binarization

$$\left\{ \begin{array}{l} \text{det. : } \begin{cases} Q(x) = 1, & x \geq 0 \\ Q(x) = -1, & \text{otherwise} \end{cases} \\ \text{sto. : } \begin{cases} P(Q(x) = 1) = \text{clip}\left(\frac{1+x}{2}, 0, 1\right) \\ P(Q(x) = -1) = 1 - P(Q(x) = 1) \end{cases} \end{array} \right. \quad \min_{\alpha \in \mathbb{R}^+, Q^B} \|\mathbf{X} - \alpha Q^B(\mathbf{X})\|_2^2, \quad \text{s.t. } Q_i^B \in \{-1, 1\}$$

- Ternarization

$$\left\{ \begin{array}{l} \text{det. : } \begin{cases} Q(x) = 1, & x > 0.5 \\ Q(x) = -1, & x < -0.5 \\ Q(x) = 0, & \text{otherwise} \end{cases} \\ \text{sto. : } \begin{cases} P(Q'(x) = 1) = \text{clip}(|x|, 0, 1) \\ P(Q'(x) = 0) = 1 - P(Q'(x) = 1) \\ Q(x) = \text{sign}(x) \cdot Q'(x) \end{cases} \end{array} \right. \quad \min_{\alpha, Q^T(\mathbf{W})} \|\mathbf{W} - \alpha Q^T(\mathbf{W})\|_2^2, \quad \text{s.t. } Q_i^T \in \{-1, 0, 1\}.$$

Advanced Quantization (k-bit quantization)

- Low-bit quantization cannot always produce acceptable accuracy
- For accuracy improvement, increase the bit-width
- Multibit quantization

$$Q^k(w) = 2Q^k \left(\frac{\tanh(w)}{2 \max(|\tanh(\mathbf{W})|)} + \frac{1}{2} \right) - 1 \quad \min_{\{\alpha_j, Q^{Bj}\}_{j=1}^k} \left\| \mathbf{X} - \sum_{j=1}^k \alpha_j Q^{Bj} \right\|_2^2, \quad \text{s.t. } Q_i^{Bj} \in \{-1, 1\}$$

Influence of Quantization

- Quantization to weight, activation and gradient is easy
- Quantization to error, update, batch normalization result in performance degradation
- Impact to model
 - $\text{Gradient} \leq \text{Weight} < \text{Activation} < \text{Error} < \text{Update}$
 - Quantization to update lower convergence
- Reduce the communication bandwidth in distributed training systems
- Gaussian distribution \rightarrow bound range of quantization loss

Network Sparsification

- Sparsification Data Object
 - **What** to be sparsified
- Sparsification Consideration
 - **How** to sparsify

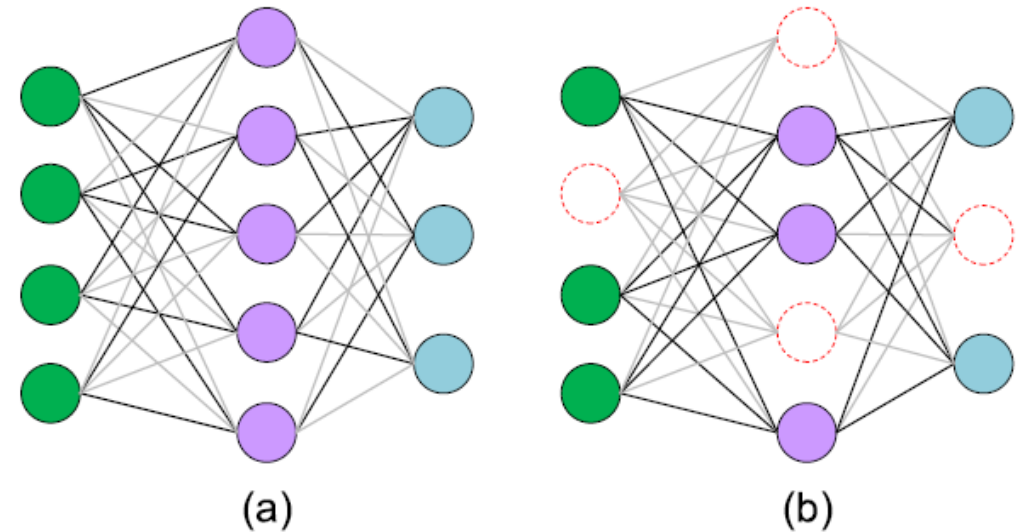


Fig. 17. Sparsification data objects. (a) Weight pruning.
(b) Neuron pruning.

Reduce #operands \rightarrow reduce #memory_access, #operations

Sparsification Data Object

- Two types of quantization target
 - Correspond to weight parameters (Weight , Gradient)
→ Weight pruning (reduce #edge)
 - Indicate neuron (Activation, Error)
→ Neuron pruning (reduce #node)
- Sparse pattern of Update is coupled with other data objects (weight, activation, error)

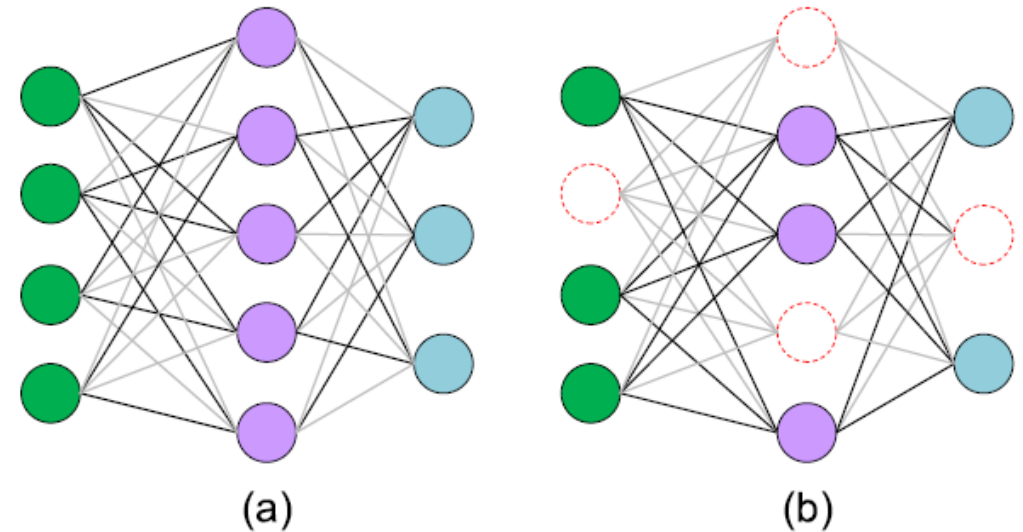


Fig. 17. Sparsification data objects. (a) Weight pruning. (b) Neuron pruning.

Sparsification Consideration

- Data Object
- Bit-width
- Sparsification grain/structures
- Different pass (forward/backward)

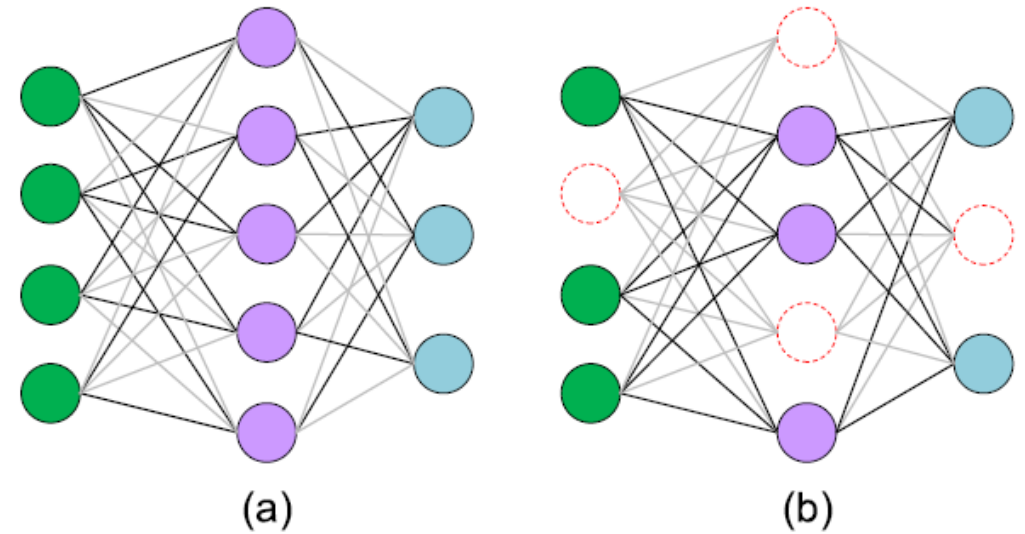


Fig. 17. *Sparsification data objects. (a) Weight pruning. (b) Neuron pruning.*

Sparsification grain/structures

- Element-wise → weight pruning
 - extra index overhead
- Vector-wise → weight pruning
- Block-wise → neuron pruning
 - Iteratively reordering the element-wise sparse weights

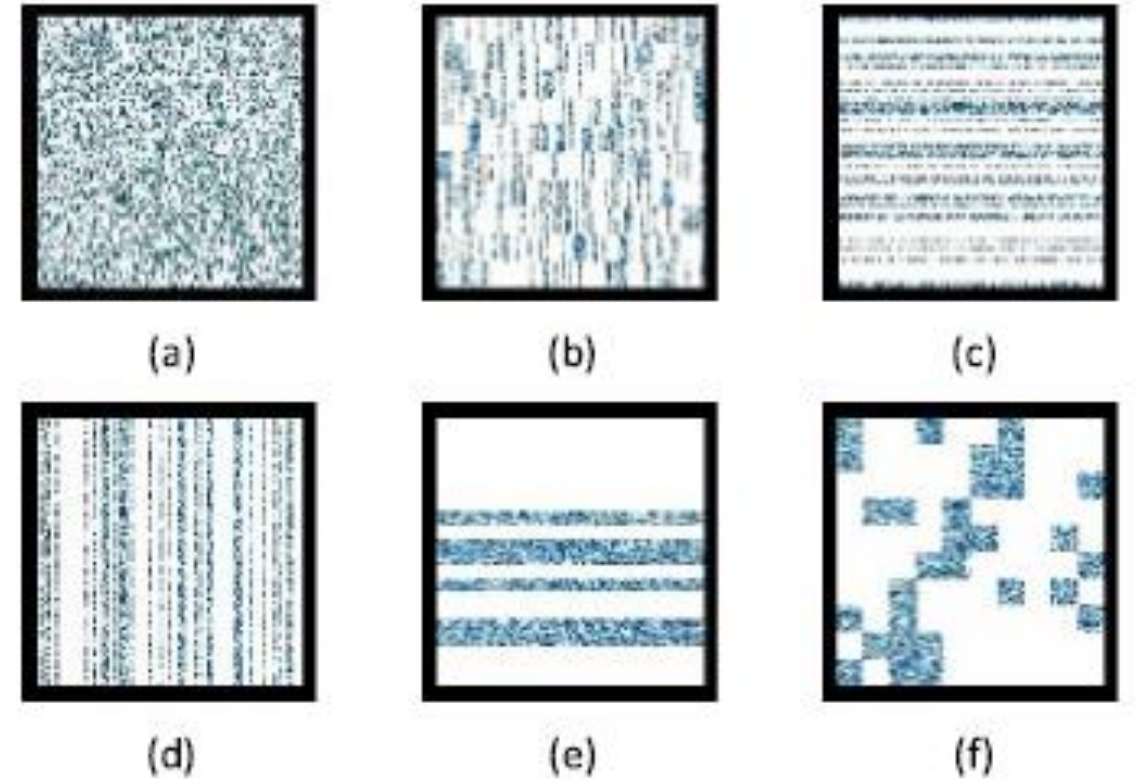
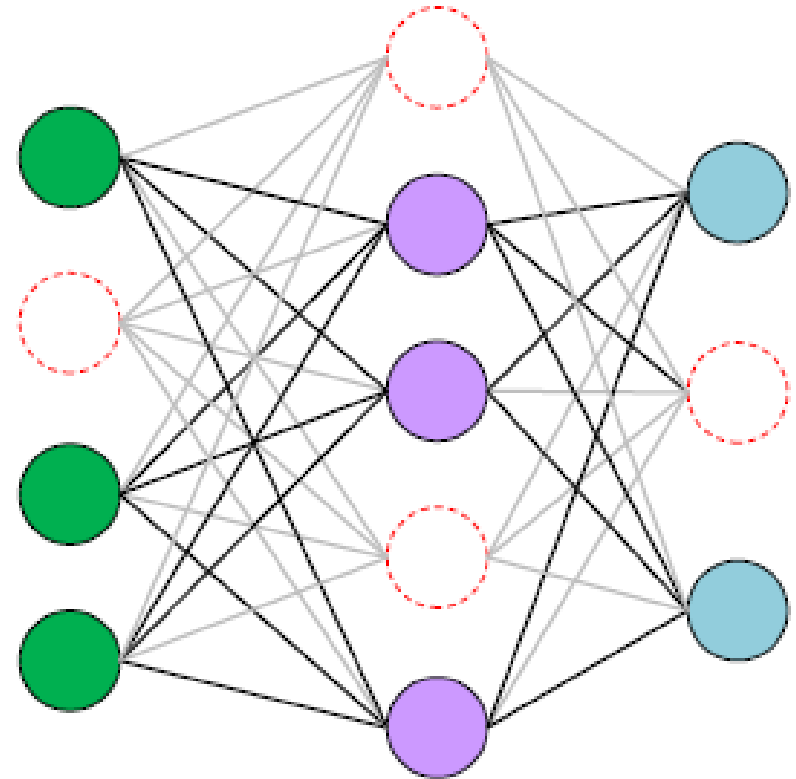


Fig. 18. Sparsification structure. (a) Element-wise. (b)–(d) Vector-wise. (e) and (f) Block-wise.

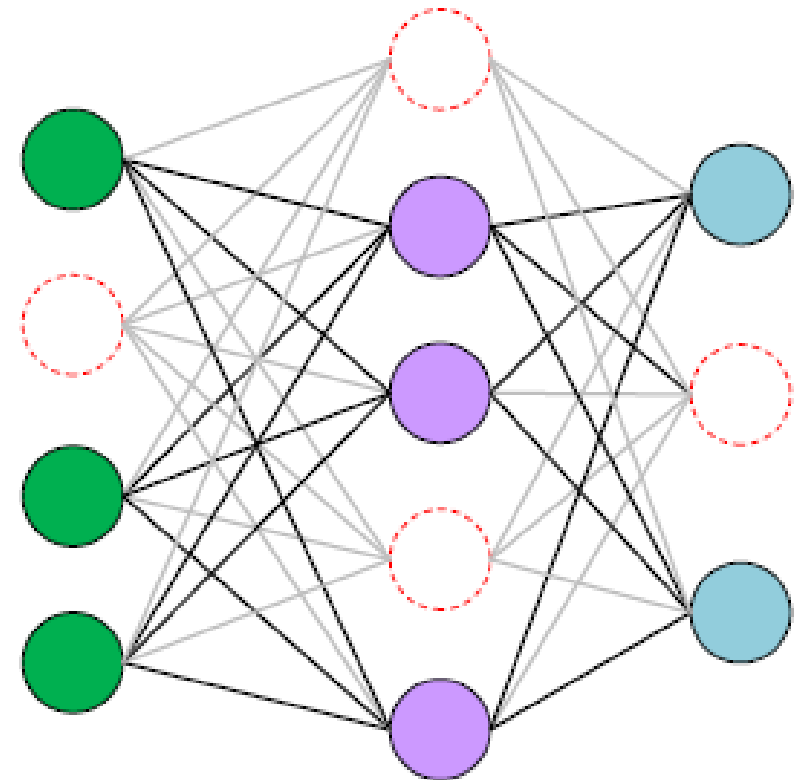
Constraints of Neuron pruning

- Target data object
 - Activation (output)
 - Error
- After pruning one input FM, these can be removed
 - Corresponding weight filter in the current layer
 - Weight channel in the next layer
- Require **complete calculation** of output activations before pruning



Before-calculation Neuron pruning

- Require **complete calculation** of output activations before pruning
- Preselection of unimportant neurons
- Hashing search
 - Negative-activation prediction
 - Low-precision estimation
 - Prechecking neighboring pixels
 - Dimension-reduction estimation



Joint-way compression

- Compact Model and Others
- Tensor Decomposition and Network Sparsification
- Data Quantization and Network Sparsification

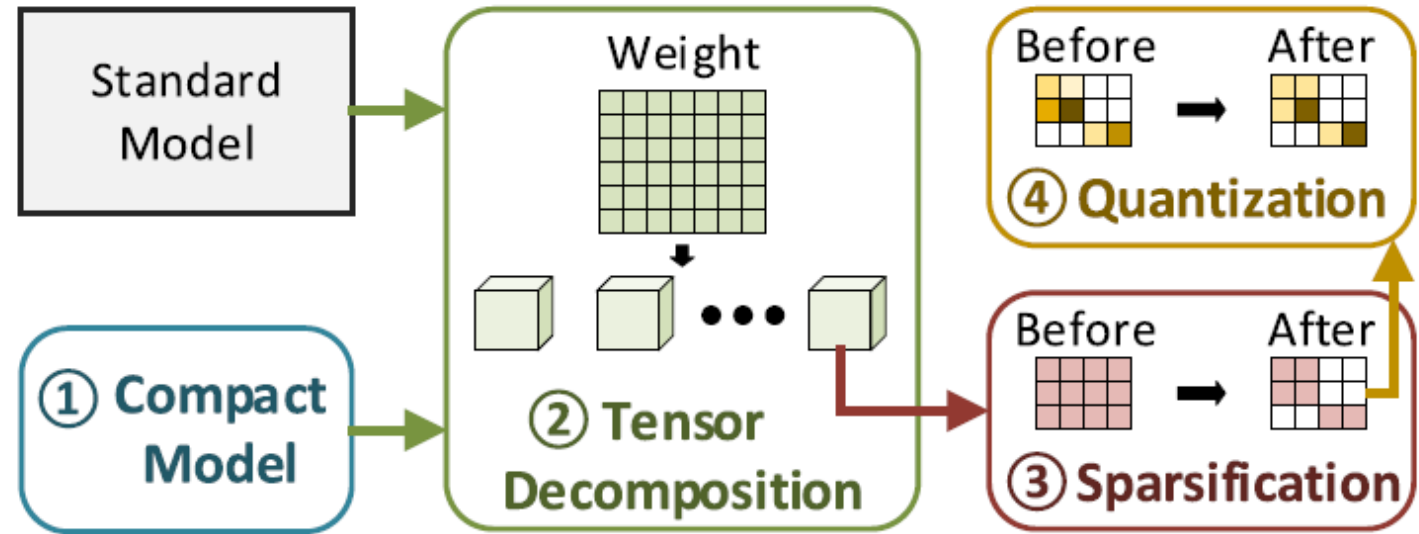


Fig. 19. *Joint-way compression.*

Joint-way compression

- Compact Model and Others
 - Compact model approach design the base network itself
 - Others compress a given base model
 - design of compact model is **orthogonal** to other compression approaches
- Tensor Decomposition and Network Sparsification
 - Sparsity by post-pruning < direct pruning of undecomposed model

Joint-way compression

- Data Quantization and Network Sparsification
- Quantization improve sparsity

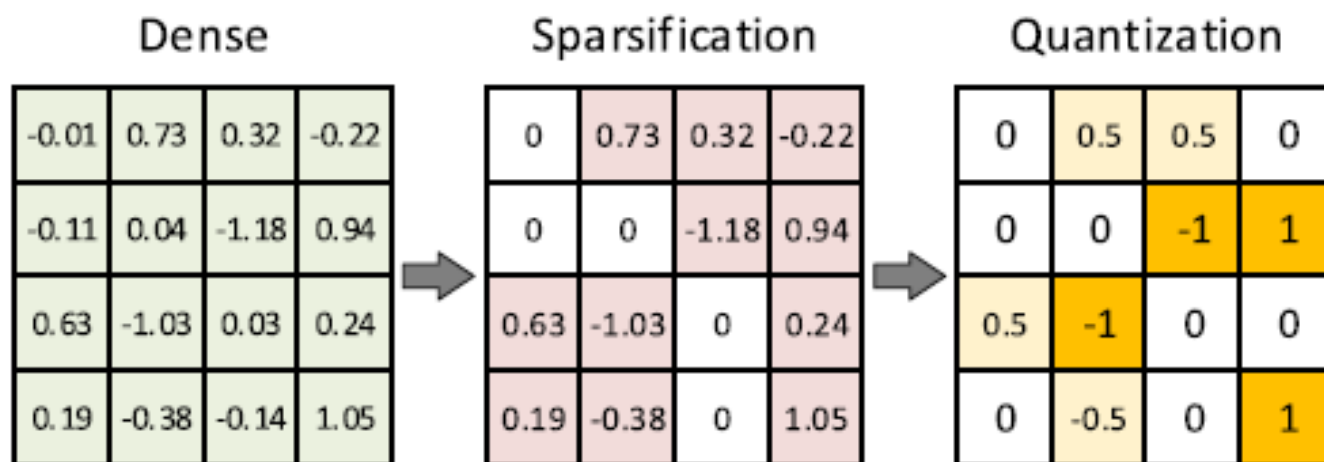


Fig. 20. *Increasing sparsity through quantization*

Neural Network Acceleration (Hardware)

- Two typical workloads in running neural networks
- Conv operation
 - Huge data reusability (activation, weights)
- MVM operation
 - Data cannot be reused without batching technique

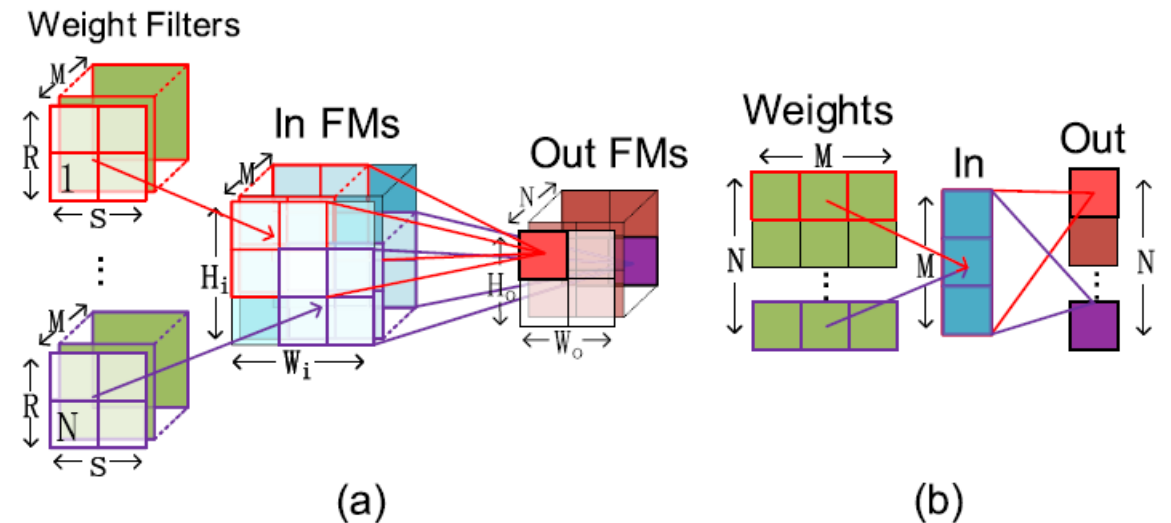


Fig. 21. Computation pattern. (a) Conv layer. (b) FC layer. Adapted from [14].

Domain-Specific Accelerator for Neural Network

- Rapid development of neural networks heavily **relies on computing capability of the hardware**
- Design goal is to minimize latency, energy, area
- On-GPU compression shows performance improvements below expectations
 - Redundant design of general-purpose processors for flexible programmability and general applicability
→ **motivation of specialized accelerator for neural network**
- Optimization targets
 - Processing architecture
 - Memory hierarchy
 - Dataflow mapping

Sole Hardware Optimization

- Parallel Compute Units and Orchestrated Memory Hierarchy
- Two aspect to accelerator design
 - Enhancing the compute parallelism
 - Optimizing the memory hierarchy

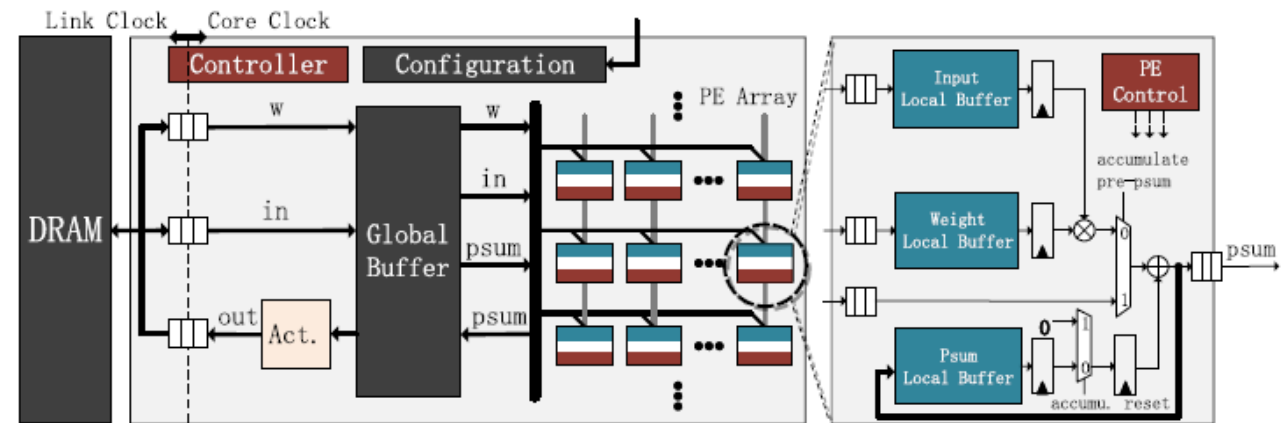


Fig. 22. Typical digital architecture of neural network accelerators. Adapted from [261].

Dataflow

- Output stationary
- Weight stationary
- Input stationary
- Row stationary

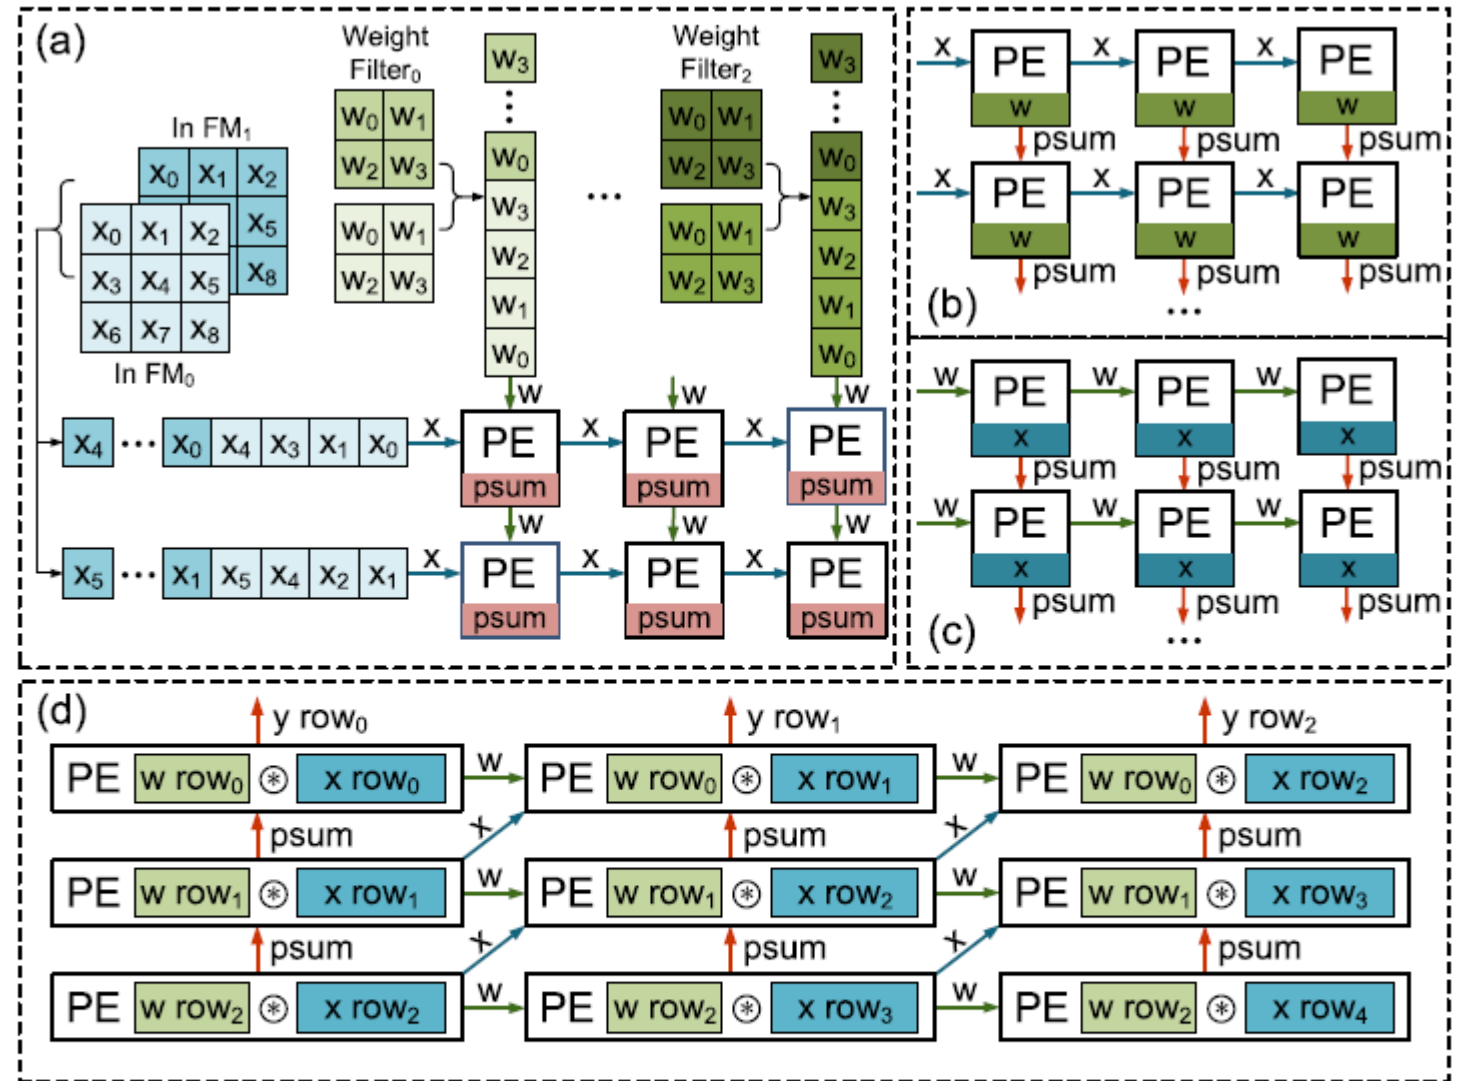


Fig. 23. Various dataflow designs. (a) Output stationary. (b) Weight stationary. (c) Input stationary. (d) Row stationary. Adapted from [14], [257], and [263].

Algorithm and Hardware Codesign

- Sole hardware side optimization meet performance wall
 - Parallelism and data reuse are exhausted
 - both algorithm and hardware side optimization
- Four Types of Codesign space
 - Transformation into Compact Model
 - Tensorized Processing Engine
 - Quantized Architecture
 - Sparse Architecture

Transformation into Compact Model

- Compact model reduce the model size without changing dense execution pattern
- Compact connections/operations and variable data path degrade the hardware utilization
- Two solutions
 - Considering flexible tiling/mapping schemes and an efficient NoC infrastructure (high bandwidth/data reuse)
 - Transform compressed network into compact models

Tensorized Processing Engine

- Decomposed network = dense matrix operation with smaller size
- Need additional support when it meet some abnormal operations included (e.g. reshaping)

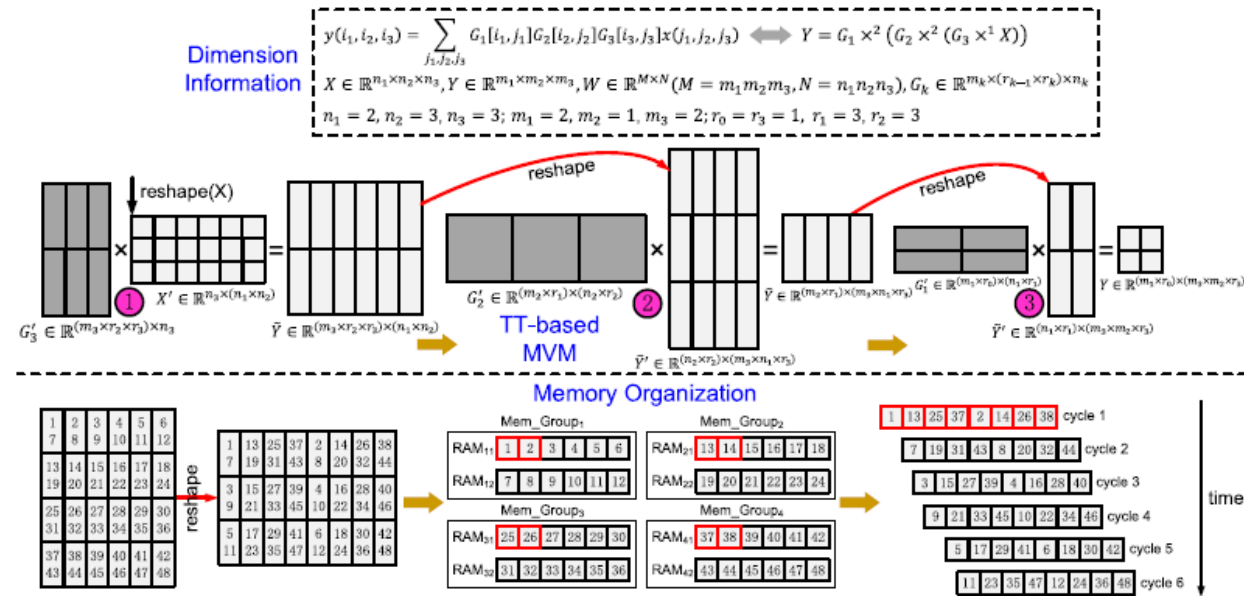


Fig. 25. TT-based MVM and an example of its memory organization. Adapted from [11].

Tensorized Processing Engine

- Accelerate TT decomposed neural network
- Reshaping
 - partitioning the working SRAM into multiple groups
 - Data selection mechanism

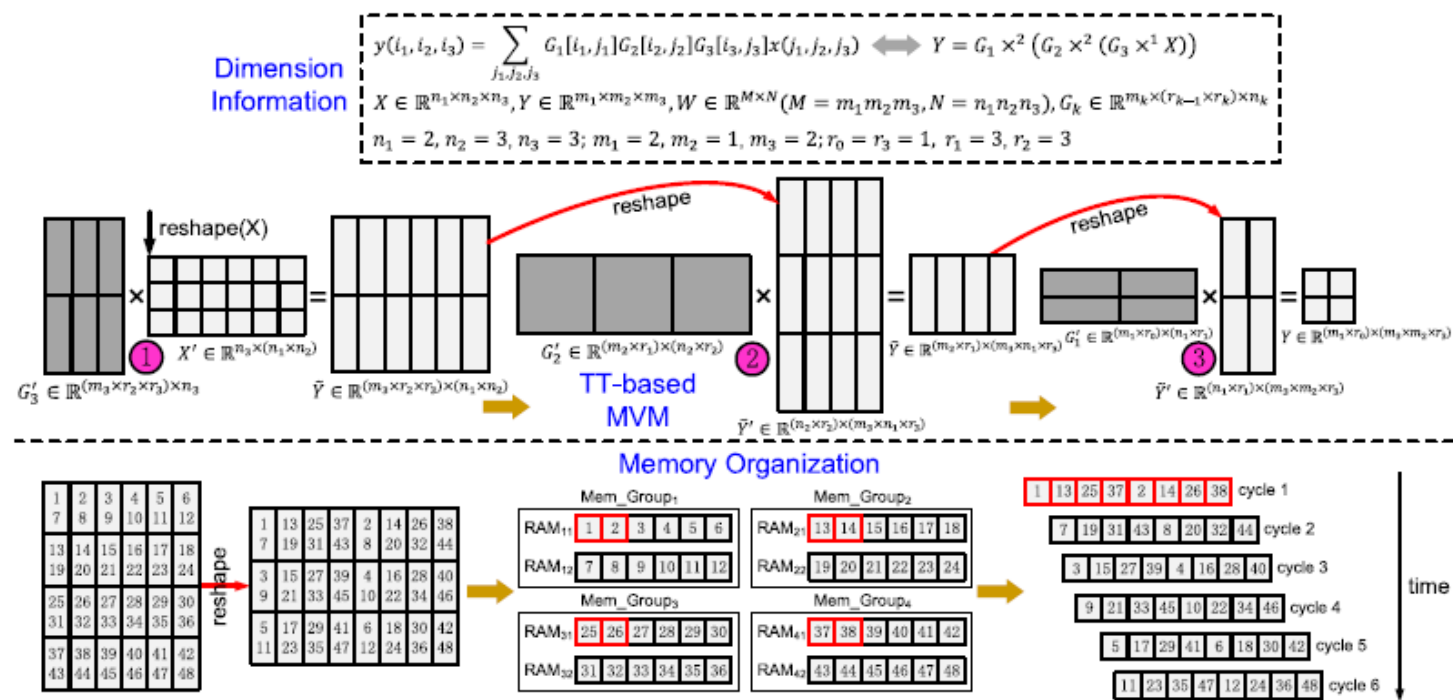


Fig. 25. TT-based MVM and an example of its memory organization. Adapted from [11].

Quantized Architecture

- Data-precision \geq 8-bit fixed point for CNN inference
- Aggressive low-bit quantization
 - Ultrahigh execution performance + accuracy loss
- Two types of quantized architecture
 - Fixed bitwidth
 - Variable bitwidth

Quantized Architecture – Fixed bitwidth

- Only change high-precision multipliers to low-bit multipliers
 - Other overall architecture and dataflow remain same
- Extreme quantization with only binary/ternary data precision
- Weight and activations are binarized/ternarized
 - MAC operation can be implemented by simple XNOR and pop-count logic operations

Quantized Architecture – Variable bitwidth

- Different layers/models require different bit-width
- Two types of architecture
 - Bit-serial
 - Bit-decomposed

Quantized Architecture – Variable bit-width

Bit-serial architecture

- Serially feed one of the two operands bit by at each cycle
- MAC operation → AND logic operation, shifted accumulation
- Enhance the parallelism

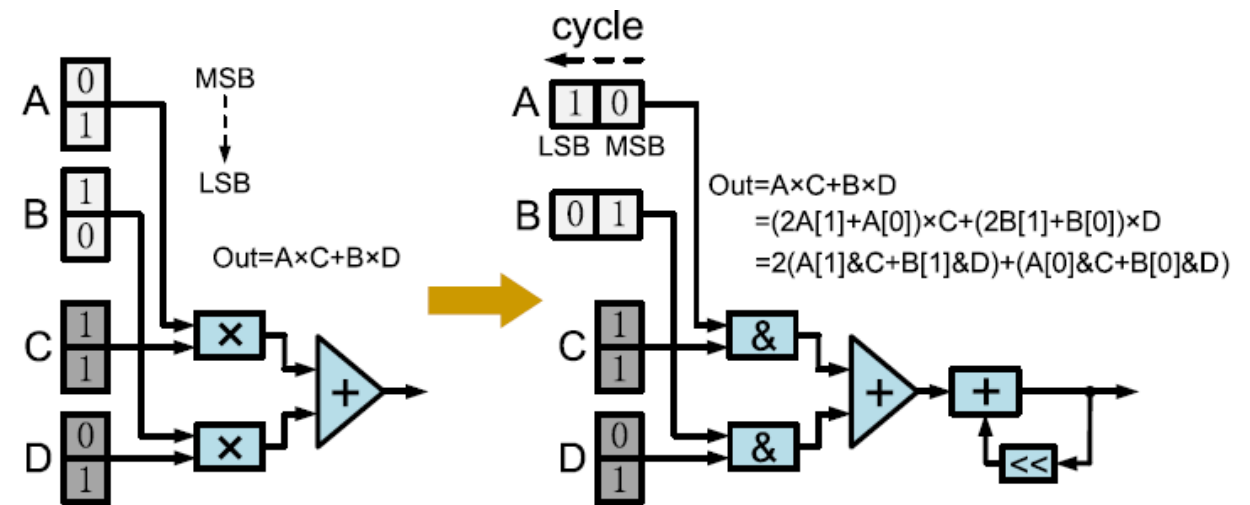


Fig. 27. Bit-serial MAC processing. Adapted from [278].

Quantized Architecture – Variable bitwidth

Bit-decomposed architecture

- Fuse multiple low-precision MACs to compute high-precision MAC
- Fusion group size is configurable to support flexible bitwidth

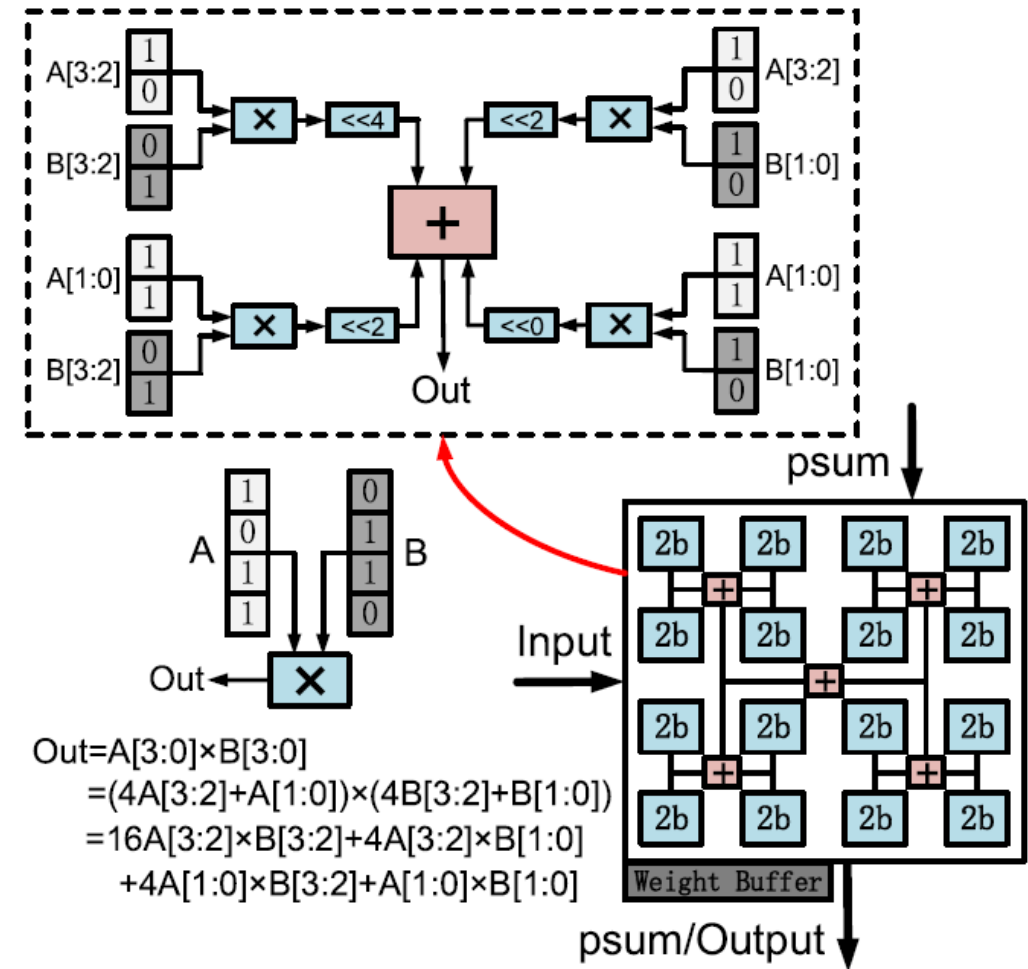


Fig. 28. Bit fusion architecture based on bit-decomposed MAC processing. Adapted from [279].

Spase Architecture

- Four exploitable sparse patterns
 - Weight sparsity
 - Input sparsity
 - Output sparsity
 - Combination of multiple patterns

Spase Architecture – Weight sparsity

- Naively skipping the MACs with zero weights
- Each PE access the required activation according to the weights index
 - Performs the residual MACs

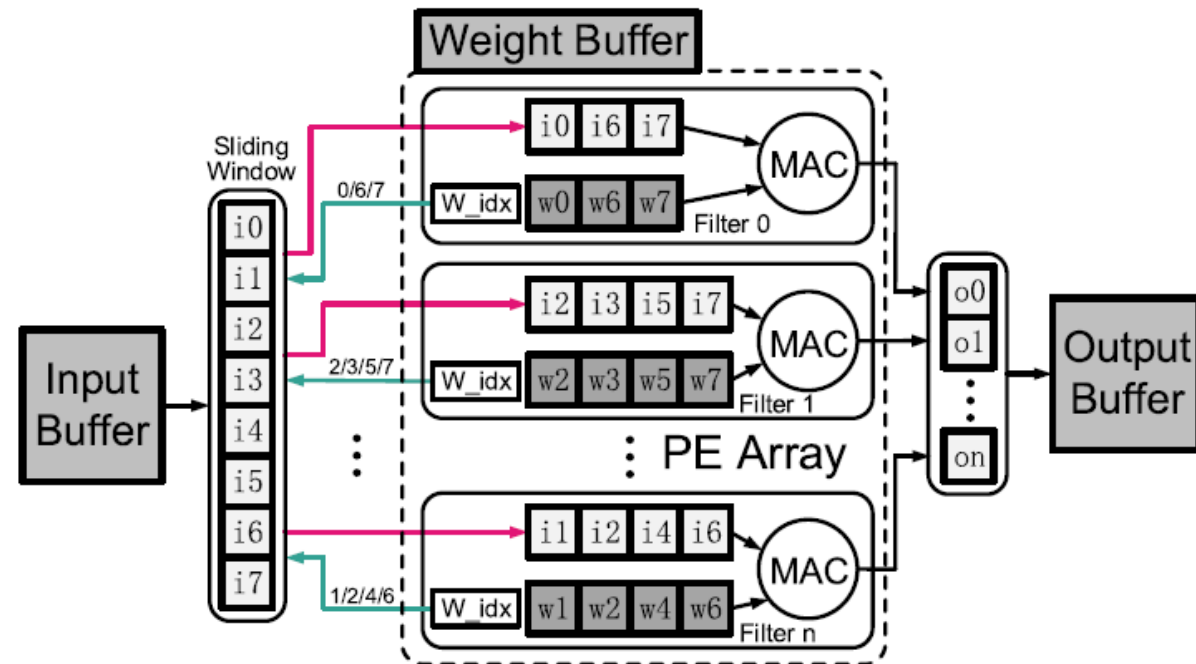


Fig. 29. Example of accelerators with weight sparsity. Adapted from [281].

Spase Architecture – Input sparsity

- To save bandwidth and energy
 - Run-length activation compression
 - Zero input detection
 - PE data gating

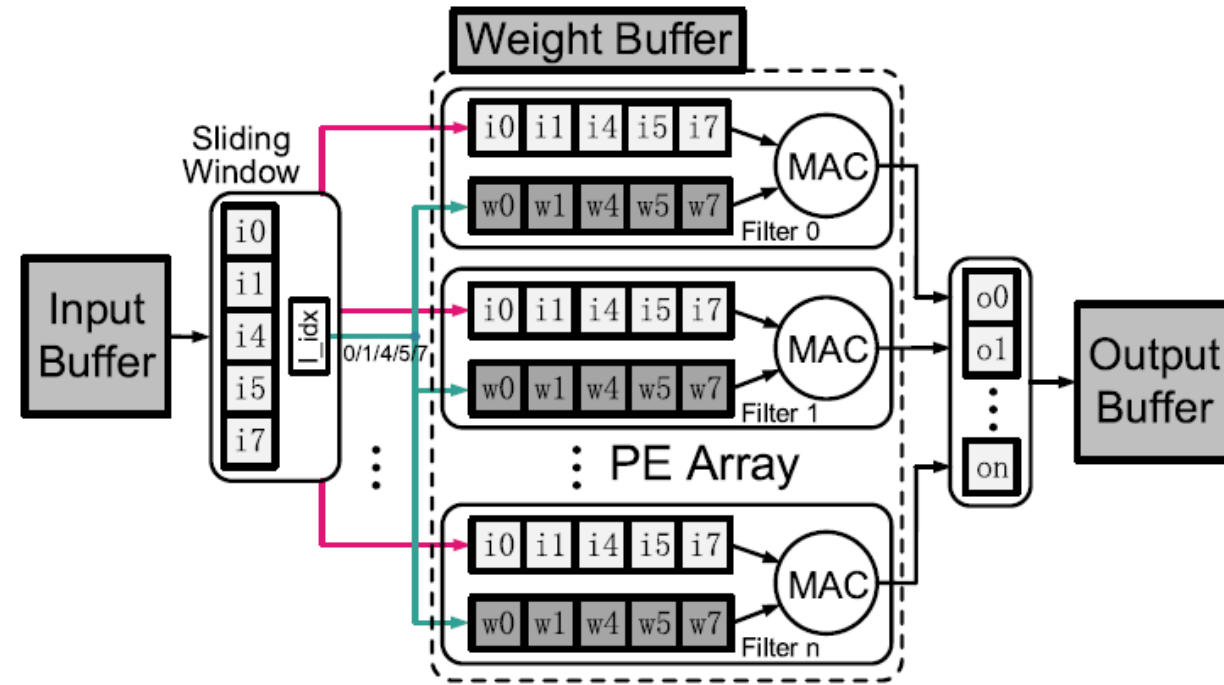


Fig. 31. Example of accelerators with input sparsity. Adapted from [283].

Spase Architecture – Output sparsity

- Predict unimportant neurons
- In-place speculation
 - Monitor the sign of the accumulated activation in each PE
 - Terminate its computation once a negative values is detected
- Decoupled speculation and execution
 - Prediction separately
 - Execute the sparse computation with predicted output index

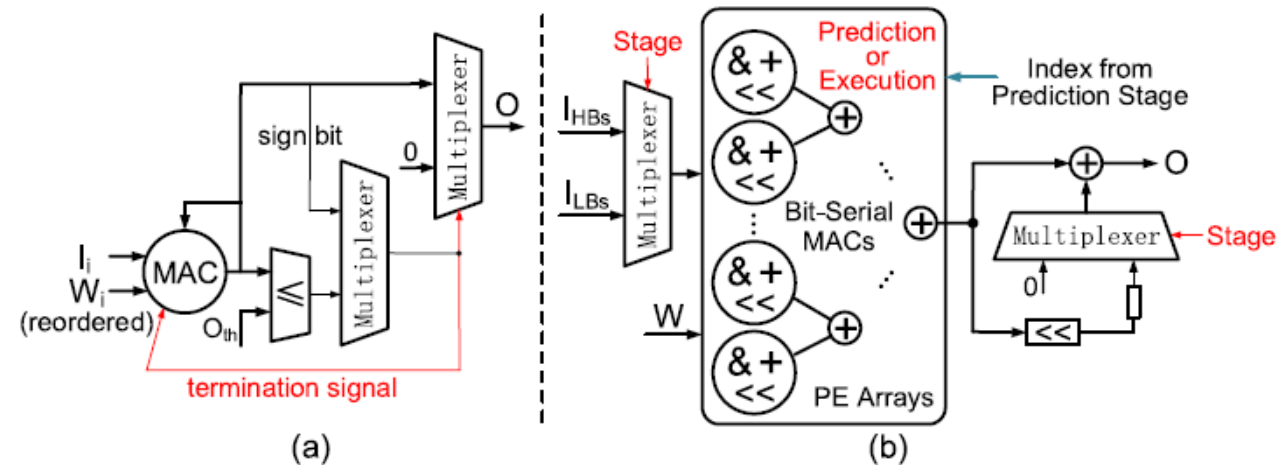


Fig. 32. Examples of accelerators with output sparsity through (a) in-place speculation or (b) decoupled speculation and execution. Adapted from [242] and [244].

Summary

- 모델 압축이 모델의 성능 저하를 일으키는 경향이 있지만, 무시할 정도의 수준으로 낮출 수 있음.
- 단일 하드웨어 최적화 측면에서 병렬화와 데이터 재사용에 대한 한계점들이 있었음
- 하드웨어와 압축 알고리즘을 함께 설계하는 방안이 증가하고 있음

Reference

- [Model compression and hardware acceleration for neural networks: A comprehensive survey.](#) Proceedings of the IEEE, 2020
- [Nvidia Technical Blog](#)

내용 정리

- [notions](#)